

Class 9

Courtney Cameron PID:A69028599

Importing candy data

```
candy_file <- read.csv("candy-data.txt", row.names=1)

head(candy_file)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer
100 Grand	1	0	1	0	0	1
3 Musketeers	1	0	0	0	1	0
One dime	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0

	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q how many chocolate candy types are in the data set?

37 of the candies are chocolate

```
length(grep('1', candy_file$chocolate))
```

```
[1] 37
```

Q how many fruity candy types are in the data set?

38 of the candies are fruity

```
length(grep('l', candy_file$fruity))
```

```
[1] 38
```

Q how many total candy?

85 different types of candies

```
nrow(candy_file)
```

```
[1] 85
```

Q how many variables/dimesnsions are there

there are 12 different variables/candy types

```
ncol(candy_file)
```

```
[1] 12
```

Data exploration using the package skimr

calling only one function from a package

```
#library(skimr)
```

```
skimr::skim(candy_file)
```

Table 1: Data summary

Name	candy__file
Number of rows	85
Number of columns	12
<hr/>	
Column type frequency: numeric	12
<hr/>	
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

What is your favorite candy

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

```
candy_file['Baby Ruth',]$winpercent
```

```
[1] 56.91455
```

Q What are the 5 least liked candies

```
inds <- order(candy_file$winpercent)
```

```
head(candy_file[inds,])
```

	chocolate	fruity	caramel	peanutyalmondy	nougat			
Nik L Nip	0	1	0		0	0		
Boston Baked Beans	0	0	0		1	0		
Chiclets	0	1	0		0	0		
Super Bubble	0	1	0		0	0		
Jawbusters	0	1	0		0	0		
Root Beer Barrels	0	0	0		0	0		
	crispedricewafer	hard	bar	pluribus	sugarpercent	pricepercent		
Nik L Nip		0	0	0	1	0.197		0.976

Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
Root Beer Barrels	0	1	0	1	0.732	0.069
	winpercent					
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					
Root Beer Barrels	29.70369					

Q4. What is the winpercent value for “Kit Kat”?

```
candy_file['Kit Kat',]$winpercent
```

[1] 76.7686

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy_file['Tootsie Roll Snack Bars',]$winpercent
```

[1] 49.6535

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The winpercent column of the data set is on a different scale than the rest of the data set

Q7. What do you think a zero and one represent for the candy\$chocolate column?


The zero means that the candy was not chosen and the 1 represents that the candy was chosen based on the randomly generated candy matchups

Q8. Plot a histogram of winpercent values

```
library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.2.3

```
ggplot(candy_file, aes(winpercent)) +  
  geom_histogram(binwidth = 8)
```



Q9. Is the distribution of winpercent values symmetrical?

the distribution is not symmetrical, it is left skewed >Q10. Is the center of the distribution above or below 50%?

the center of the distribution is below 50% >Q11. On average is chocolate candy higher or lower ranked than fruit candy?

-find all chocolate candy (subset) -get the winpercent values -summarize these values into one metric (mean)

```
choc <- subset(candy_file$winpercent, candy_file$chocolate==1)  
mean(choc)
```

```
[1] 60.92153
```

```
fruity <- subset(candy_file$winpercent, candy_file$fruity==1)  
mean(fruity)
```

```
[1] 44.11974
```

```
t.test(choc,fruity)
```

Welch Two Sample t-test

```
data: choc and fruity
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

Q12. Is this difference statistically significant?

the difference is statistically significant with a p-value of 2.87e-08

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
inds <- order(candy_file$winpercent)

head(candy_file[inds,])
```

	chocolate	fruity	caramel	peanut	almond	nougat		
Nik L Nip	0	1	0		0	0		
Boston Baked Beans	0	0	0		1	0		
Chiclets	0	1	0		0	0		
Super Bubble	0	1	0		0	0		
Jawbusters	0	1	0		0	0		
Root Beer Barrels	0	0	0		0	0		
	crisped	rice	wafer	hard	bar	pluribus	sugar	percent
Nik L Nip		0	0	0		1	0.197	0.976
Boston Baked Beans		0	0	0		1	0.313	0.511
Chiclets		0	0	0		1	0.046	0.325
Super Bubble		0	0	0		0	0.162	0.116

Jawbusters	0	1	0	1	0.093	0.511
Root Beer Barrels	0	1	0	1	0.732	0.069

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744
Root Beer Barrels	29.70369

Q14. What are the top 5 all time favorite candy types out of this set?

```
inds <- order(candy_file$winpercent)

tail(candy_file[inds,])
```


	chocolate	fruity	caramel	peanut	almond	nougat
Reese's pieces	1	0	0		1	0
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0

	crisped	rice	wafer	hard	bar	pluribus	sugar
Reese's pieces		0	0	0		1	0.406
Snickers		0	0	1		0	0.546
Kit Kat		1	0	1		0	0.313
Twix		1	0	1		0	0.546
Reese's Miniatures		0	0	0		0	0.034
Reese's Peanut Butter cup		0	0	0		0	0.720

	price	percent	winpercent
Reese's pieces	0.651	73.43499	
Snickers	0.651	76.67378	
Kit Kat	0.511	76.76860	
Twix	0.906	81.64291	
Reese's Miniatures	0.279	81.86626	
Reese's Peanut Butter cup	0.651	84.18029	

Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy_file, aes(winpercent, rownames(candy_file))) +
  geom_col()
```




Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by winpercent?

creating color vector to color the graph based on the type of candies

```
my_cols=rep('darkgrey', nrow(candy_file))
my_cols[as.logical(candy_file$chocolate)] = "chocolate4"
my_cols[as.logical(candy_file$bar)] = "brown"
my_cols[as.logical(candy_file$fruity)] = "pink3"
```

plot was saved because the aspect ratio was bad for readability

```
ggplot(candy_file, aes(winpercent, reorder(rownames(candy_file), winpercent))) +
  geom_col(fill=my_cols) +
  labs(x='Matchup Win Percent', y=NULL)
```

```
ggsave('barplot.png', height=10, width=7)
```

how to insert any image caption text can go in the square brackets adding a {} bracket with #fig, a hyperlink will be added with a link to the graph

Q17. What is the worst ranked chocolate candy?

Sixlets

Q18. What is the best ranked fruity candy?

Nik L Nip

Taking a look at pricepoint

```
ggplot(candy_file, aes(x=winpercent, y=pricepercent, label=rownames(candy_file)))+
  geom_point(col=my_cols) +
  geom_text(col=my_cols)
```





Figure 1: win percentage




ggrepel package can be used to make point labels look nicer, geom_text will put in the labels but doesn't account for overlap

```
library(ggrepel)
```

Warning: package 'ggrepel' was built under R version 4.2.3

```
ggplot(candy_file, aes(x=winpercent, y=pricepercent, label=rownames(candy_file)))+
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, max.overlaps = 8)
```

Warning: ggrepel: 61 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reese's miniatures

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

The least popular is the Nik L Nip


```
inds2 <- order(candy_file$pricepercent, decreasing=TRUE)
head(candy_file[inds2,c(11,12)], n=5)
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Q21. Make a barplot again with `geom_col()` this time using `pricepercent` and then improve this step by step, first ordering the x-axis by value and finally making a


so called “dot chat” or “lollipop” chart by swapping `geom_col()` for `geom_point()` + `geom_segment()`.

```
ggplot(candy_file) +
  aes(pricepercent, reorder(rownames(candy_file), pricepercent)) +
  geom_segment(aes(yend = reorder(rownames(candy_file), pricepercent),
                  xend = 0), col="gray40") +
  geom_point()
```



```
ggsave('barplot2.png', height = 10,width = 7)
```

Because aspect ratio makes seeing the graph difficult, the plot is loaded in as an image



Exploring Correlation

```
library(corrplot)
```

Warning: package 'corrplot' was built under R version 4.2.2


corrplot 0.92 loaded

```
cij <- cor(candy_file)
```

```
head(cij)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat
chocolate	1.0000000	-0.7417211	0.24987535	0.37782357	0.25489183
fruity	-0.7417211	1.0000000	-0.33548538	-0.39928014	-0.26936712
caramel	0.2498753	-0.3354854	1.00000000	0.05935614	0.32849280
peanutyalmondy	0.3778236	-0.3992801	0.05935614	1.00000000	0.21311310
nougat	0.2548918	-0.2693671	0.32849280	0.21311310	1.00000000
crispedricewafer	0.3412098	-0.2693671	0.21311310	-0.01764631	-0.08974359
	crispedricewafer	hard	bar	pluribus	sugarpercent
chocolate	0.34120978	-0.3441769	0.5974211	-0.3396752	0.10416906
fruity	-0.26936712	0.3906775	-0.5150656	0.2997252	-0.03439296
caramel	0.21311310	-0.1223551	0.3339600	-0.2695850	0.22193335
peanutyalmondy	-0.01764631	-0.2055566	0.2604196	-0.2061093	0.08788927
nougat	-0.08974359	-0.1386750	0.5229764	-0.3103388	0.12308135
crispedricewafer	1.00000000	-0.1386750	0.4237509	-0.2246934	0.06994969
	pricepercent	winpercent			
chocolate	0.5046754	0.6365167			
fruity	-0.4309685	-0.3809381			
caramel	0.2543271	0.2134163			
peanutyalmondy	0.3091532	0.4061922			
nougat	0.1531964	0.1993753			
crispedricewafer	0.3282654	0.3246797			

```
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and fruity are negatively correlated

Q23. Similarly, what two variables are most positively correlated?

chocolate highly correlates with bar and winpercent

PCA

PCA needs to be scaled - scale=TRUE

```
pca <- prcomp(candy_file, scale=TRUE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369


	PC8	PC9	PC10	PC11	PC12
--	-----	-----	------	------	------

Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

```
pca_result <- as.data.frame(pca$x)
```

```
ggplot(pca_result, aes(x=PC1,y=PC2,label=rownames(pca_result)))+
  geom_point(col=my_cols)+
  geom_text_repel(col=my_cols, max.overlaps = 7, size=3)
```


Warning: ggrepel: 51 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Loadings plot

```
pca.loadings <- as.data.frame(pca$rotation)
```

```
ggplot(pca.loadings, aes(PC1, reorder(rownames(pca.loadings),PC1)))+
  geom_col()
```



Q24. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you?

The positive direction mostly picked up the fruity, pluribus, and hard candies. This makes sense because these three characteristics are commonly found together within fruity flavored candies